

Modern Approaches to the Rich Vehicle Routing Problem

Henry F. R. Fellows
Division of Science and Mathematics
University of Minnesota, Morris
Morris, Minnesota, USA 56267
fello056@morris.umn.edu

ABSTRACT

The Rich Vehicle Routing Problem is a class of problems that revolve around finding the most optimal route for a certain set of deliveries. Obtaining approximate solutions to these computationally hard problems is both economically valuable and [Henry: academically interesting?](#). Despite a long history of research in the field, new methods and approaches to specific subproblems are frequently introduced. This paper summarizes the problem, its applications, recent algorithms and approaches.

Keywords

Vehicle Routing Problem, Black-box optimization, Genetic Algorithms, Probability Collectives, Distributed Systems, [Henry: & more - user interfaces](#),

1. INTRODUCTION

[Henry: Routing is cool and important, also Amazon prime, Uber, self driving cars, and mail.](#)

Routing is a deeply pervasive feature of the modern world. From the point-to-point navigation of everyday driving to the shipping of bulk freight, the task of creating those routes has been increasingly automated. The long term trend towards self-driving vehicles means that computers will soon control the entire process of transportation; aside from selecting a departure and destination, humans need no longer be involved. Software to create routes for school buses, mail deliveries and garbage trucks have been in use for decades; everything from Amazon to Uber uses these systems on a daily basis. The problem of finding efficient routes quickly is still an ongoing problem in the real world. The exact impact of improvements in routing is difficult to measure, but in 2011, the external costs of traffic jams in the EU were 1-2% of GDP[2]. A small improvement - say five percent - in routing efficiency represents a staggering amount of money. The general case of routing is the traveling salesman problem. The traveling salesman problem asks for the shortest route which passes through each point once. Assuming that




Figure 1: Example of the VRP
[Henry: Needs to be replaced w/own image.](#)

each pair of points is connected by a link, the number of potential routes is $0.5n!$, which grows extremely quickly, and is significant even for small values. For $n = 6$, the number of routes reaches 20,160. The traveling salesman problem is in a class of problems known as NP-complete. Although it has not yet been proven, it is considered likely that there is no algorithm for NP-complete problems that allows them to be solved quickly. Fortunately, solutions of NP-complete problems are easy to verify. In practice, the traveling salesman problem is too abstract for real-world routing problems, and so a variant was proposed that better represents the challenges of routing vehicles.

1.1 Routing in the real world

Routing in the real world is often a complicated and messy affair. The vehicle has to have enough capacity for all of its deliveries, and often, customers expect that the vehicle arrives near a specified time. New stops may be added during the day, or the deliveries may not be known ahead of time. In some cases, the customer might not know how much of a product they need until the vehicle arrives. Dealing with these types of problems is the domain of the Rich Vehicle Routing Problem.

2. THE VEHICLE ROUTING PROBLEM

Originally titled the Truck Dispatching Problem, the original formulation of the Vehicle Routing Problem (VRP) was

This work is licensed under the Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> or send a letter to Creative Commons, 171 Second Street, Suite 300, San Francisco, California, 94105, USA.

UMM CSci Senior Seminar Conference, December 2015 Morris, MN.

created by G. B. Danzig and J. H. Ramser in 1959[3]. The premise of the problem is that each vehicle (or truck), has a limited capacity and the fleet must make deliveries to as many customers as possible, starting from a specific location known as a depot[2]. Alternatively, the vehicles may have to make multiple deliveries in order to satisfy all of the customers. A measure that is also considered worth minimizing is the maximum tour length - the longest single route taken by one truck. Danzig and Ramser note that if the total capacity of the vehicles is less than the total demand of the customers, the problem is mathematically identical to the traveling salesman problem. There has been a great deal of research done on the original VRP, and recent research focuses on versions of the VRP with different constraints or multiple constraints simultaneously. These variations of the Vehicle Routing Problem are collectively referred to as the Rich Vehicle Routing Problem[2] (RVRP). In the following sections of, we discuss prominent variants of the RVRP.

2.1 Decentralized Vehicle Routing Problem

Henry: It's agent flavored! Many problems in the rich vehicle routing category feature decentralized control[2], where each vehicle or depot is modeled as having an independent agent who makes decisions out of self-interest. However, pure self-interest can cause decisions that degrade the solutions of others. Imagine a vehicle which finds a route that fulfills a large number of deliveries extremely quickly; other vehicles might find that their routes force them to cross through this area where they cannot make any deliveries along the way. While it is highly efficient for the first vehicle, it is less efficient for all other vehicles. The ideal is to pursue optimal routes - both on the local (individual) scale, and on a global scale. Finding other methods of making decisions that allow for both autonomy and global utility is an ongoing research question.

2.2 Vehicle Routing Problem with Time Windows

Henry: Deliveries during business hours only! From the corner of the divan of Persian saddle-bags on which he was lying, smoking, as was his custom, Innumerable cigarettes, Lord Henry Wotton could just catch the gleam of the honey-sweet and honey-coloured blossoms of a laburnum, whose tremulous branches seemed hardly able to bear the burden of a beauty so flame-like as theirs ; and now and then the fantastic shadows of birds in flight flitted across the long tussore-silk curtains that were stretched in front of the huge window, producing a kind of momentary Japanese effect, and making him think of those pallid jade-faced painters of Tokio who, through the medium of an art that is necessarily immobile, seek to convey the sense of swiftness and motion. The sullen murmur of the bees shouldering their way through the long unmown grass, or circling with monotonous insistence round the dusty gilt horns of the straggling woodbine, seemed to make the stillness more oppressive.

2.3 Approaches

In general the approaches used to solve the VRP are approximate methods; they intend to find 'good', not perfect solutions. The algorithms tend to use simple rules to guess solutions, and the process of narrowing down or perfecting these guesses is known as optimization. Henry: History, maybe?. Most modern methods of solving the VRP

belong to the blackbox style of optimization; blackbox optimization is used when a problem does not have a formal algebraic model, or the model is too computationally expensive [?]. The VRP does have a very good model, but it is extremely computationally expensive to use it, and so the natural choice is to use blackbox methods. The common feature of blackbox methods is the use of stochastic (random) elements, especially in selecting starting states. This makes blackbox methods into a double edged sword; they are often effective at solving otherwise intractable problems, but the solutions frequently fail to produce any insight into the problem. Stochastic decision making makes it hard to determine why the algorithm provides a specific solution.

Henry: Maybe: The preference for blackbox methods may be a reflection of the difficulty in finding/improving [non-guessing] algorithms (or, perhaps equivalently, of the inherent computational limitations of calculating a perfect answer to the VRP)

3. GENETIC AND MEMETIC ALGORITHMS

Henry: Explain local search Henry: Of Memetic Algorithms: In general, the genetic algorithm improves the solution in large strokes, while the local optimization fine tunes the solutions generated by the GA. GAs are cool too, and the HGSADC is a pretty spiffy algorithm From the corner of the divan of Persian saddle-bags on which he was lying, smoking, as was his custom, Innumerable cigarettes, Lord Henry Wotton could just catch the gleam of the honey-sweet and honey-coloured blossoms of a laburnum, whose tremulous branches seemed hardly able to bear the burden of a beauty so flame-like as theirs ; and now and then the fantastic shadows of birds in flight flitted across the long tussore-silk curtains that were stretched in front of the huge window, producing a kind of momentary Japanese effect, and making him think of those pallid jade-faced painters of Tokio who, through the medium of an art that is necessarily immobile, seek to convey the sense of swiftness and motion. The sullen murmur of the bees shouldering their way through the long unmown grass, or circling with monotonous insistence round the dusty gilt horns of the straggling woodbine, seemed to make the stillness more oppressive.

From the corner of the divan of Persian saddle-bags on which he was lying, smoking, as was his custom, Innumerable cigarettes, Lord Henry Wotton could just catch the gleam of the honey-sweet and honey-coloured blossoms of a laburnum, whose tremulous branches seemed hardly able to bear the burden of a beauty so flame-like as theirs ; and now and then the fantastic shadows of birds in flight flitted across the long tussore-silk curtains that were stretched in front of the huge window, producing a kind of momentary Japanese effect, and making him think of those pallid jade-faced painters of Tokio who, through the medium of an art that is necessarily immobile, seek to convey the sense of swiftness and motion. The sullen murmur of the bees shouldering their way through the long unmown grass, or circling with monotonous insistence round the dusty gilt horns of the straggling woodbine, seemed to make the stillness more oppressive.

3.1 Human Assisted Routing

Henry: Move here? Uses a GA as base algorithm, is a memetic algorithm uses local search

Humans tend to be very good at solving visual problems,




Figure 2: To be example of selecting good/bad regions

Henry: Needs to be replaced w/own image.

and this extends to the routing if the problem is displayed appropriately. We have a talent for deciding whether a given route is good or bad by simply looking at a diagram. Amazon’s mechanical turk and similar platforms were created to allow developers to make use of human intelligence in solving computationally complex problems. Pursuing the approach of using humans to do things humans are good at anyways, S. Ismail, F. Legras and G. Coppin[1] created a genetic algorithm that uses humans to determine how good a particular solution is.

3.1.1 Interactive Genetic Algorithm for the VRP

```

Initialize population;
while number of interactions without improvement
  <  $lt_{NI}$ , and time <  $T_{max}$  do
    Genetic algorithm go! Human selection go! Good
    tagging - higher selection probability - mutation
    operator not allowed Bad tagging - lower selection
    probability - mutation operator encouraged
    Standard heuristic function
  end
return best feasible solution;

```

3.2 Hybrid Genetic Search with Advanced Diversity Control

Adding time window constraints to customer demand and depot availability poses a significant challenge, especially in relation to the smaller proportion of solutions and the increased computational complexity caused by the increased search space. The opposing demands of temporal and spatial constraints is a particularly frustrating problem in the VRPTW.

The Hybrid Genetic Search with Advanced Diversity Control (HGSADC) address some of these challenges in the

VRPTW, especially in route-duration constraints. The primary feature of the HGSADC is the different approach to diversity management in its population. HGSADC adds diversity to its objectives as a term to be optimized, which allows it to avoid dead-end solutions. The algorithm is considered the current state of the art in the multi-depot vehicle routing problem with time windows[7].

3.2.1 Algorithm and mechanics

HGSADC is a complex algorithm, and cannot be fully described in this paper. What follows is a summary of the algorithm and the most notable features.

```

Initialize population;
while number of interactions without improvement
  <  $lt_{NI}$ , and time <  $T_{max}$  do
    Select parent solutions  $P_1$  and  $P_2$ ;
    Create offspring  $C$  from  $P_1$  and  $P_2$  (crossover);
    Educate  $C$  (local search procedure);
    if  $C$  infeasible then
      Insert  $C$  into infeasible subpopulation;
      Repair with probability  $P_{rep}$ ;
    end
    if  $C$  feasible then
      Insert  $C$  into feasible subpopulation;
    end
    if maximum subpopulation size reached then
      Select survivors;
    end
    if best solution not improved for  $It_{div}$  iterations
      then
      Diversify population;
    end
    Adjust penalty parameters for infeasibility;
    if number of iterations =  $k \times It_{dec}$  where  $k \in \mathbb{N}$ 
      then
        Henry: k made up of natural numbers?;
        Decompose the master problem;
        Use HGSADC on each subproblem;
        Reconstitute three solutions, and insert them in
        the population;
      end
    end
  end
return best feasible solution;

```

Henry: Algorithm-figure goes here eventually. HGSDAC evolves infeasible and feasible solutions as two separate subpopulations. Two parents are selected by genetic operators Henry: ??? iteratively applied to the population, and are then combined into an offspring. A local search is then applied to the offspring, and then may be repaired Henry: explain or skip and placed into an appropriate subpopulation. If a subpopulation exceeds a maximum size, a survivor selection stage is triggered Henry: explain.... A diversification round may be started if there has been It_{div} iterations without improvement. Henry: to be cont, pending decisions on what to gloss over.

3.3 Adaptive Memetic Algorithms

Henry: Section may be removed

4. AGENT-BASED MODELS

Agents are simply small decision making elements that typically have some internal state and the ability to make decisions based on this internal state. In optimization, a typical goal is for an agent to attempt to achieve some sort of optimal state. Agents (should) take actions to improve their state, and we describe this as the agent acting in its own interest. The function that returns the quantification of this local optimization is termed the local utility function, in contrast to the global, or system utility function. The global utility function describes how optimal the system is, which can be quite different from the local utility. **Henry: Make up definition? Agents are very poorly defined. “A set of agents, with a function $f(x)$ that change their internal state, based on external state. They may change external state as well.”** We have chosen to discuss two major approaches in modeling individuals for the decentralized VRP. The first is agent-based models, the more traditional approach, and the second is probability collectives, a relatively recent approach that has a legacy in genetic algorithms.

4.1 Distributed reverse Vickrey auctions

Saleh et al.[5] were examining a version of the DCVRP where each depot and customer are treated as agents. The authors were interested in designing a system that encouraged near-optimal solutions while still being completely self-interested. They pursued this by building mechanisms that encouraged depots to offer an accurate estimate of minimum costs while still allowing the cost functions of each depot to remain private. **Henry: reverse vickrey auction**

4.1.1 The Distributed RVA Routing Algorithm

There is a set D of m agent depots; each depot, k , has N_k vehicles with total capacity q_k^i . Each vehicle is denoted $i_k^1 \dots i_k^{N_k}$, and assigned to route r_k^i . The number of vehicles, capacity, and the cost function are private to each depot. The set of customers is C , and a single customer is z , which has a demand d_z that can be served only by one vehicle. This information is public to all depots.

For every customer j , we estimate the additional cost that results from adding j to the customers served by depot k , for each depot k such that $q_k^i \geq d_z$. The least additional cost dc_k is obtained using an approximate solution.

Then, find the depot k with the lowest dc_k , and add the customer to r_k , the route indicated by the method of estimating the least additional cost. The routes of depot k can be optimized by applying an exact solution to the VRP to the order of the list of nodes in each route. Finally, the capacity of depot k is updated to the sum of $q_k^i - d_j$.

Henry: Alternatively:

Each depot k submits a bid to each unassigned customer j where $q_k^i \geq d_z$. Bids consist of the estimated costs of servicing the customer. The customer chooses the lowest bid satisfying their demands and notifies k^* . The depot k^* may receive responses from other customers. From the given responses, the depot chooses the customer j^* with the lowest insertion cost (the customer whose addition to the vehicle's route will cause the least change in the routing cost). Once the customer is assigned to a depot, the customer j^* submits payment to depot k^* . The payment is equal to the second lowest offer of the game. This, combined with the fact that bids are sealed, encourages depots to provide the lowest possible rate (because if they win, they are still paid more than the estimate). The completion of these steps is

a single auction round, and many rounds may take place until every customer is assigned. Once the auction rounds are complete, the depot can apply other optimization strategies to its routes, or in the case of small routes, an exact solution.

4.2 Probability Collectives

Probability Collectives (PC) takes a different approach to selecting a good solution. Instead of attempting to evolve an exemplary individual in a population in the method of genetic or memetic algorithms, probability collectives selects optimal strategy for each agent[4]. A PC agent is a self-interested, learning individual that selects a strategy with the highest probability of optimizing the local and global objectives. In contrast to the Distributed Reverse Vickrey Auction, the agents willingly and freely share information about their strategies, and the cost function is global.

4.2.1 Mechanics of Probability Collectives

The general form of PC in the context of VRP is as follows, along with the flowchart below. Imagine you have N vehicles, each having a set X of strategies, of (potentially variable) length m . The strategies are most often routes in RVRP applications[6], but can sometime be as large as entire sets of routes or as small as individual actions. The set is represented as $X_i = \{X_i^{[1]}, X_i^{[2]}, \dots, X_i^{[m]}\}$.

A strategy for vehicle i is denoted as $X_i^{[r]}$, where r is an identifier for that strategy. For each vehicle, assign a uniform probability, $1/m_i$ to all actions; the resulting probability of that strategy is signified by $q(X_i^{[r]})$. The vehicle then selects a random action r and a sampling of random actions from other vehicles. The resulting set, the ‘combined strategy set’, $Y_i^{[r]} = \{X_1^{[?]}, \dots, X_i^{[r]}, \dots, X_N^{[?]}\}$, represents a guess as to a potential future. Accordingly, for each of the strategy sets $Y_i^{[r]}$, compute the expected local utility - the measure of how good the solution is expected to be for the agent - using the following measure:

$$\text{Expected Utility of Agent } i^r = q_i^r \prod_{(i)} q(x_{(i)}^{[?]}) \cdot G(Y_i^{[r]}) \quad (1)$$

Where q_i^r represents the probability of action r for vehicle i , (i) is the set of all agents excluding i , and G is the function that computes global utility for a given set of strategies. G is problem specific, but in the RVRP it could be a measure of unused capacity, unvisited destinations, maximum tour length or various combinations of measurements.

After computing the local utility, the next step is to update the probability of each action for all vehicles as follows:

$$q(X_i^{[r]}) < -q(X_i^{[r]}) - \alpha_{\text{Step}} \cdot q(X_i^{[r]}) \cdot k_{r \text{ update}} \quad (2)$$

where k is the iteration, and

$$k_{r \text{ update}} = \frac{\text{contrib. of Agent } i}{T} + S_i(q) + \ln(q(X_i^{[r]})) \quad (3)$$

Here, Boltzmann's temperature T is a scalar that represents the importance of the of the contribution of agent i . **Henry: This is only the literal interpretation of T . It might also be $T(k) = T_0/\ln(k)$, but I can't find any good reference.** The contribution of Agent i is then:

$$\text{contrib. Agent } i = \text{exp. utility Agent } i^r - \text{exp. global utility} \quad (4)$$

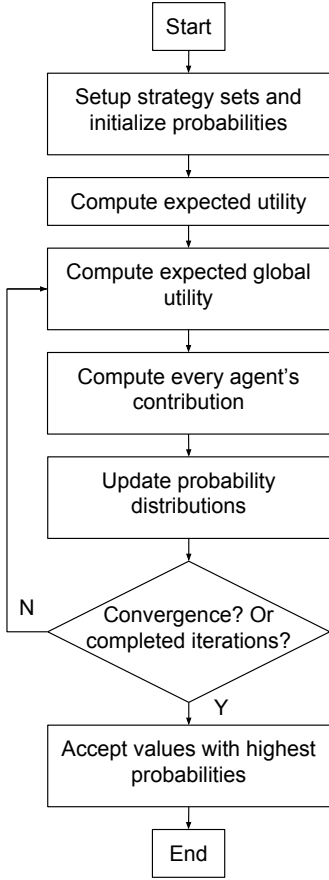


Figure 3: Probability Collectives Algorithm.
 Henry: THIS FIGURE IS GIVING ME AN ULCER.

The $S_i(q)$ term is the entropy of every agent’s probability distribution. Entropy in the context of information theory is the expected value of the information in a message. As it increases, the probability distribution more clearly distinguishes the contribution of each strategy toward optimizing the expected global utility. Entropy is computed by:

$$S_i(q) = - \sum_{r=1}^{m_i} q(X_i^{[r]}) \cdot \ln(q(X_i^{[r]})) \quad (5)$$

To summarize, if a strategy r creates a larger contribution to the optimization of the objective than other strategies, the probability associated with r increases by a larger amount. This entire process is repeated until the probability distribution converges or the maximum number of iterations are completed. The strategy with the highest probability is then returned.

5. CONCLUSION

From the corner of the divan of Persian saddle-bags on which he was lying, smoking, as was his custom, Innumerable cigarettes, Lord Henry Wotton could just catch the gleam of the honey-sweet and honey-coloured blossoms of a laburnum, whose tremulous branches seemed hardly able to bear the burden of a beauty so flame-like as theirs ; and

now and then the fantastic shadows of birds in flight flitted across the long tussore-silk curtains that were stretched in front of the huge window, producing a kind of momentary Japanese effect, and making him think of those pallid jade-faced painters of Tokio who, through the medium of an art that is necessarily immobile, seek to convey the sense of swiftness and motion. The sullen murmur of the bees shouldering their way through the long unmown grass, or circling with monotonous insistence round the dusty gilt horns of the straggling woodbine, seemed to make the stillness more oppressive.

From the corner of the divan of Persian saddle-bags on which he was lying, smoking, as was his custom, Innumerable cigarettes, Lord Henry Wotton could just catch the gleam of the honey-sweet and honey-coloured blossoms of a laburnum, whose tremulous branches seemed hardly able to bear the burden of a beauty so flame-like as theirs ; and now and then the fantastic shadows of birds in flight flitted across the long tussore-silk curtains that were stretched in front of the huge window, producing a kind of momentary Japanese effect, and making him think of those pallid jade-faced painters of Tokio who, through the medium of an art that is necessarily immobile, seek to convey the sense of swiftness and motion. The sullen murmur of the bees shouldering their way through the long unmown grass, or circling with monotonous insistence round the dusty gilt horns of the straggling woodbine, seemed to make the stillness more oppressive.

6. ACKNOWLEDGEMENTS

I’d like to thank caffeine, water, and stress; the raw elements that formed this paper. Harris L. Mayer’s 1964 paper, “Opacity Calculations, Past and Future” was a wonderful source of thematic inspiration.

7. REFERENCES

- [1] S. Amaran, N. V. Sahinidis, B. Sharda, and S. J. Bury. Simulation optimization: a review of algorithms and applications. *4OR*, 12(4):301–333, 2014.
- [2] S. Ben Ismail, F. Legras, and G. Coppin. A new interactive evolutionary algorithm for the vehicle routing problem. In *Proceedings of the 14th Annual Conference Companion on Genetic and Evolutionary Computation, GECCO ’12*, pages 661–668, New York, NY, USA, 2012. ACM. *Fascinating idea that explores a best-of-both worlds approach.*
- [3] J. Caceres-Cruz, P. Arias, D. Guimarans, D. Riera, and A. A. Juan. Rich vehicle routing problem: Survey. *ACM Comput. Surv.*, 47(2):32:1–32:28, Dec. 2014. *This paper is /the/ summary of the field at the moment. It supports the background and describes the major techinques in use in the field.*
- [4] G. B. Dantzig and J. H. Ramser. The truck dispatching problem. *Management Science*, 6(1):80–91, 1959. *This paper presents the original problem and context.*
- [5] A. J. Kulkarni and K. Tai. Probability collectives for decentralized, distributed optimization: A collective intelligence approach. In *2008 IEEE International Conference on Systems, Man and Cybernetics*, pages 1271–1275, Oct 2008.
- [6] M. Saleh, A. Soeanu, S. Ray, M. Debbabi, J. Berger, and A. Boukhtouta. Mechanism design for

decentralized vehicle routing problem. In *Proceedings of the 27th Annual ACM Symposium on Applied Computing*, SAC '12, pages 749–754, New York, NY, USA, 2012. ACM. *decentralization*.

- [7] M. Vasirani and S. Ossowski. Decentralized coordination strategies for the vehicle routing problem. In *Proceedings of the 2008 ACM Symposium on Applied Computing*, SAC '08, pages 130–131, New York, NY, USA, 2008. ACM. *Unique focus on total decentralization*.
- [8] T. Vidal, T. G. Crainic, M. Gendreau, and C. Prins. A hybrid genetic algorithm with adaptive diversity management for a large class of vehicle routing problems with time-windows. *Computers and Operations Research*, 40(1):475 – 489, 2013.